REFERENCE · R2

Task Delegation Decision Guide

AI Literacy for the Modern Workforce

Use this guide before delegating a task to AI. Walk through four questions in sequence. Each question either routes you to a recommendation or sends you to the next question.

Q1

STAKES

What happens if this output is wrong?

CONSIDER: *Who sees this? What decisions depend on it? What's the cost of an error: rework, reputational damage, legal exposure, financial loss?*

MINIMAL CONSEQUENCE	MODERATE TO SIGNIFICANT
FULLY DELEGABLE	Continue to **Q2** ↓

Q2

CONTEXTUAL KNOWLEDGE

Does this task require knowledge the model doesn't have?

CONSIDER: *Does it depend on unrecorded conversations, relationship context, proprietary data you haven't provided in the prompt, or real-time information the model can't access?*

YES	NO
HUMAN-ONLY	Continue to **Q3** ↓

Q3

VERIFIABILITY

Can you verify every key claim in the output against an independent source?

CONSIDER: *Is there a dashboard, database, document, or known standard you can check against? Or would you be relying on the output's plausibility alone? If you will run this task many times, also ask whether the check scales: can you verify a hundred outputs as cheaply as one, or does each one need a person? A check that needs a person each time limits how many outputs you can delegate; it does not on its own make the answer No. One Anthropic engineer put it this way: delegate where validation effort is small compared with creation effort.*

NO	YES
HUMAN-ONLY	Continue to **Q4** ↓

Q4

DOMAIN EXPERTISE

Do you have the expertise to catch errors in this output?

CONSIDER: *Would you recognize a wrong number, a flawed argument, or a misapplied concept in this domain? Or would a plausible-sounding error pass your review?*

NO	YES
HUMAN-ONLY	**AI-ASSISTED WITH REVIEW**

Note: If Q4 = No but the task still benefits from AI acceleration, use AI-assisted with mandatory peer review from someone with the domain expertise. Without that expertise, errors are hard to spot. In a survey of 319 knowledge workers, 58 reported barriers to inspecting AI responses, such as lacking the domain knowledge to judge them, and some who doubted their own ability accepted outputs by default. Anthropic's engineers are divided: some prefer familiar ground where they can verify the output, others use AI in peripheral domains to save time, and one engineer, when unsure of a topic, asks the model for options to research rather than answers to accept.

RECOMMENDATION REFERENCE

FULLY DELEGABLE

WHAT IT MEANS

AI produces the output from a clear task definition; you spot-check the result.

TYPICAL TASKS

Data compilation from known sources, formatting, proofreading, standard translations, template population.

AI-ASSISTED WITH REVIEW

WHAT IT MEANS

AI produces a draft or structure; you direct, evaluate, and revise using your expertise.

TYPICAL TASKS

Comparative analysis with verified inputs, first-draft summaries you'll review, structural scaffolding.

HUMAN-ONLY

WHAT IT MEANS

You produce the output. AI may assist with sub-components, but the core judgment stays with you.

TYPICAL TASKS

Strategic recommendations, causal interpretation, tasks depending on unrecorded or proprietary information.

ONE TASK, MULTIPLE COMPONENTS

Most workplace tasks are bundles of sub-tasks with different risk profiles. Apply this guide *per component*, not per deliverable. A quarterly business review might include data compilation (fully delegable), trend narrative (AI-assisted), and strategic recommendations (human-only). The framework names this split directly: deconstructing a task into AI, human, and collaborative components.

When you decompose a task into separate components, a compilation step followed by a dedicated human verification step can be fully delegated even when the overall deliverable carries significant stakes, because the verification step absorbs the risk. If compilation and verification are not separated, the combined task routes through the full flowchart at the higher stakes level.

When Anthropic interviewed 53 of its own engineers and researchers about what they hand to Claude, three of this guide's four tests turned up in their answers: stakes, verifiability, and context. They delegated work where quality is not critical and where checking the output costs less than producing it. They kept work that needs context the model lacks, along with strategic and design judgment. They also favored self-contained components, the per-component rule above. The expertise test is where they split: some keep AI to familiar ground they can verify, while others hand it low-complexity work outside their own experience. In the earlier survey of 132 staff, more than half said they could fully delegate only a fifth of their work or less. Source: Huang et al., Dec 2025[1] (Anthropic engineers; findings may not generalize).

BEFORE THE FOUR QUESTIONS: THE DATA GATE

The flowchart assumes the task is yours to delegate. One gate comes first: whatever you paste into an AI tool leaves your organization's boundary the moment you send it. MIT's 2026 AI committee notes that most commercial AI plans, especially paid ones, let users switch off the use of their data for training. The default, however, is often to record sessions for training. Whether client records, unreleased financials, personal data, internal strategy, or anything under NDA can enter an external tool at all is an organizational policy decision, not an individual delegation decision. The framework covers the tool side of this inside Delegation: Platform Awareness includes evaluating AI tools against operational and regulatory needs. The material side stays your organization's call.

Check your organization's AI-use policy before you start the flowchart. If no policy exists, that absence is worth raising. R7's Delegation Norms section gives you a conversation starter for this question.

Framework: 4D AI Fluency Framework, Delegation competency (Dakan & Feller, 2025)[2], as taught in Anthropic's AI Fluency course[3]. Data-gate defaults: MIT Ad Hoc Committee, 2026[4]. Barriers to inspecting AI output: Lee et al., CHI 2025[5]. Engineer delegation practices: Huang et al., Dec 2025[6].

Notes

1. Saffron Huang et al., "How AI Is Transforming Work at Anthropic," Anthropic, December 2025, https://anthropic.com/research/how-ai-is-transforming-work-at-anthropic/
2. Rick Dakan and Joseph Feller, "Framework for AI Fluency: Practical Overview Document (Version 1.1)," Ringling College of Art and Design; University College Cork, January 2025, pp. 3–4, https://ringling.libguides.com/ai/framework
3. Rick Dakan, Joseph Feller, and Anthropic, "AI Fluency: Framework & Foundations (course)," Anthropic Academy, 2025, https://anthropic.skilljar.com/ai-fluency-framework-foundations
4. MIT Ad Hoc Committee on AI Use in Teaching, Learning, and Research Training, "Report of MIT's Ad Hoc Committee on AI Use in Teaching, Learning, and Research Training," Massachusetts Institute of Technology, August 2026, pp. 25–26, https://sites.mit.edu/ai-use/files/2026/08/AI-Committee-Final-Report-Aug-13.pdf
5. Hao-Ping (Hank) Lee et al., "The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects From a Survey of Knowledge Workers," Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems, 2025, p. 11 (awareness, motivation, and ability barriers; trivial task 55/319; trust 83/319; time 44/319; barriers to inspecting responses, such as limited domain knowledge, 58/319; barriers to revising queries and improving output 72/319), https://doi.org/10.1145/3706598.3713778
6. Saffron Huang et al., "How AI Is Transforming Work at Anthropic," Anthropic, December 2025, research post, "AI delegation approaches": six criteria from 53 interviews, https://anthropic.com/research/how-ai-is-transforming-work-at-anthropic/

© Michael Ritchot. Code licensed under the MIT License. Instructional content licensed under CC BY-NC-SA 4.0.