Output Verification Checklist

AI Literacy for the Modern Workforce

Use this checklist before any AI-assisted output is shared, submitted, or published. Apply the same review standard you would to work from a junior colleague. Analyst Benedict Evans's shorthand is that AI gives you infinite interns (Evans, May 2026, slide 60): tasks easy to explain to an intern (Evans, May 2025, slide 27), done at any volume, with an intern's error rate. A security engineer at Anthropic made the same comparison, describing a model proposal as smart in the dangerous way a talented junior engineer can be, the kind only experience catches (Huang et al., Dec 2025).

PRODUCT DISCERNMENT

Does the output meet the specification?

- [ ] Factual claims verified against an independent source (not just plausible-sounding)
- [ ] Citations checked for existence (search for the source; don't trust surface features like journal names, authors, or page numbers)
- [ ] Statistics and specific numbers checked against original data
- [ ] Names, dates, and institutional references confirmed
- [ ] Scope matches what was requested (nothing critical added or omitted)
- [ ] Framing and defaults checked for skew the training data may have introduced (who is represented, what is treated as the norm, which perspectives are missing)
- [ ] Format and length match specification

PROCESS DISCERNMENT

Is the reasoning sound?

- [ ] Logical reasoning traced step by step, with conclusions that follow rather than just sound plausible
- [ ] Conclusions follow from stated premises (not restated as evidence)
- [ ] No unsupported assumptions smuggled in as facts
- [ ] No circular reasoning
- [ ] Dependency chains checked: if conclusions rely on upstream claims, verify upstream first
- [ ] Alternatives or counterarguments considered where appropriate

PERFORMANCE DISCERNMENT

Does the behavior match?

- [ ] Tone and register appropriate for the intended audience
- [ ] Level of directness matches what was requested
- [ ] No hedging language where direct language was specified (or vice versa)
- [ ] Voice consistent throughout (no register shifts mid-document)

VERIFICATION DECISION FRAMEWORK

When evaluating each element, classify it:

Fabricated	Generated content presenting as retrieved. *Specific citations, quotes, statistics from unfamiliar sources.*	**DEFAULT ACTION** Remove or replace. Verify against primary source.
Uncertain	Could be real or generated; you can't tell from the output. *Specific metrics, operational figures, quantitative claims.*	**DEFAULT ACTION** Verify against primary source before circulating.
Reliable-with-caveats	Manages its own attribution and limitations transparently. *Claims with explicit sourcing, noted limitations.*	**DEFAULT ACTION** Standard editorial review.
Dependency risk	Conclusions rely on upstream inputs that may be unreliable. *Recommendations, action items, strategic conclusions.*	**DEFAULT ACTION** Verify premises before trusting conclusions.

DIAGNOSTIC PAIRS: WHAT TO CHECK AND WHY

Name the mechanical pair to know what to check and how:

PAIR	PRODUCES · CHECK BY...
Next-token prediction × Knowledge	Hallucinated specifics (citations, quotes, statistics) **CHECK BY...** Search for the source independently
Context window × Steerability	Drift in long conversations (instructions forgotten, tone shifts) **CHECK BY...** Compare output against original instructions
Tokenization × Next-token prediction	Confident arithmetic errors **CHECK BY...** Recalculate independently

Framework: 4D AI Fluency Framework, Discernment competency (Dakan & Feller, 2025, pp. 5–6), as taught in Anthropic's AI Fluency course.

© Michael Ritchot. Code licensed under the MIT License. Instructional content licensed under CC BY-NC-SA 4.0.