Output Verification Checklist

AI Literacy for the Modern Workforce

Use before any AI-assisted output is shared, submitted, or published. Apply the same review standard you would to work from a junior colleague. The industry shorthand is that AI gives you infinite interns (Evans, May 2026, slide 60). Tasks easy to explain to an intern, done at any volume, with an intern's error rate. The comparison is not loose: a security engineer at Anthropic described a model proposal as smart in the dangerous way a talented junior engineer can be, the kind only experience catches (Huang et al., Dec 2025).

PRODUCT DISCERNMENT

Does the output meet the specification?

- [ ] Factual claims verified against an independent source (not just plausible-sounding)
- [ ] Citations checked for existence (search for the source; don't trust surface features like journal names, authors, or page numbers)
- [ ] Statistics and specific numbers checked against original data
- [ ] Names, dates, and institutional references confirmed
- [ ] Scope matches what was requested (nothing critical added or omitted)
- [ ] Framing and defaults checked for skew the training data may have introduced (who is represented, what is treated as the norm, which perspectives are missing)
- [ ] Format and length match specification

PROCESS DISCERNMENT

Is the reasoning sound?

- [ ] Logical reasoning traced step by step, with conclusions that follow rather than just sound plausible
- [ ] Conclusions follow from stated premises (not restated as evidence)
- [ ] No unsupported assumptions smuggled in as facts
- [ ] No circular reasoning
- [ ] Dependency chains checked: if conclusions rely on upstream claims, verify upstream first
- [ ] Alternatives or counterarguments considered where appropriate

PERFORMANCE DISCERNMENT

Does the behavior match?

- [ ] Tone and register appropriate for the intended audience
- [ ] Level of directness matches what was requested
- [ ] No hedging language where direct language was specified (or vice versa)

- [ ] Voice consistent throughout (no register shifts mid-document)

VERIFICATION DECISION FRAMEWORK

When evaluating each element, classify it:

Fabricated	Generated content presenting as retrieved. *Specific citations, quotes, statistics from unfamiliar sources.*	**DEFAULT ACTION** Remove or replace. Verify against primary source.
Uncertain	Could be real or generated; you can't tell from the output. *Specific metrics, operational figures, quantitative claims.*	**DEFAULT ACTION** Verify against primary source before circulating.
Reliable-with-caveats	Manages its own attribution and limitations transparently. *Claims with explicit sourcing, noted limitations.*	**DEFAULT ACTION** Standard editorial review.
Dependency risk	Conclusions rely on upstream inputs that may be unreliable. *Recommendations, action items, strategic conclusions.*	**DEFAULT ACTION** Verify premises before trusting conclusions.

DIAGNOSTIC PAIRS: WHAT TO CHECK AND WHY

Name the mechanical pair to know what to check and how:

PAIR	PRODUCES · CHECK BY...
NTP × Knowledge	Hallucinated specifics (citations, quotes, statistics) **CHECK BY...** Search for the source independently
Context window × Steerability	Drift in long conversations (instructions forgotten, tone shifts) **CHECK BY...** Compare output against original instructions
Tokenization × NTP	Confident arithmetic errors **CHECK BY...** Recalculate independently

Framework: 4D AI Fluency Framework, Discernment competency (Dakan & Feller, 2025, pp. 5–6), as taught in Anthropic's AI Fluency course.

© Michael Ritchot. Code licensed under the MIT License. Instructional content licensed under CC BY-NC-SA 4.0.