REFERENCE CARD · R4 DISCERNMENT

Output Verification Checklist

AI Literacy for the Modern Workforce

ONE-PAGE REFERENCE
PRINT OR SAVE

BEFORE YOU SHIP
— pause & verify

Use before any AI-assisted output is shared, submitted, or published.
Apply the same review standard you would to work from a junior colleague.

THREE CHECKS · DISCERNMENT

01 PRODUCT DISCERNMENT

Does the output meet the specification?

- [ ] Factual claims verified against an independent source (not just plausible-sounding)
- [ ] Citations checked for existence (search for the source; don't trust surface features like journal names, authors, or page numbers)
- [ ] Statistics and specific numbers checked against original data
- [ ] Names, dates, and institutional references confirmed
- [ ] Scope matches what was requested (nothing critical added or omitted)
- [ ] Format and length match specification

02 PROCESS DISCERNMENT

Is the reasoning sound?

- [ ] Logical reasoning traced step by step, with conclusions that follow rather than just sound plausible
- [ ] Conclusions follow from stated premises (not restated as evidence)
- [ ] No unsupported assumptions smuggled in as facts
- [ ] No circular reasoning
- [ ] Dependency chains checked: if conclusions rely on upstream claims, verify upstream first
- [ ] Alternatives or counterarguments considered where appropriate

03 PERFORMANCE DISCERNMENT

Does the behavior match?

- [ ] Tone and register appropriate for the intended audience
- [ ] Level of directness matches what was requested
- [ ] No hedging language where direct language was specified (or vice versa)
- [ ] Voice consistent throughout (no register shifts mid-document)

VERIFICATION DECISION FRAMEWORK

When evaluating each element, classify it:

CATEGORY	WHAT IT MEANS	DEFAULT ACTION
Fabricated	Generated content presenting as retrieved. Specific citations, quotes, statistics from unfamiliar sources.	Remove or replace. Verify against primary source.
Uncertain	Could be real or generated; you can't tell from the output. Specific metrics, operational figures, quantitative claims.	Verify against primary source before circulating.
Reliable-with-caveats	Manages its own attribution and limitations transparently. Claims with explicit sourcing, noted limitations.	Standard editorial review.
Dependency risk	Conclusions rely on upstream inputs that may be unreliable. Recommendations, action items, strategic conclusions.	Verify premises before trusting conclusions.

DIAGNOSTIC PAIRS: WHAT TO CHECK AND WHY

Name the mechanical pair to know what to check and how:

PAIR	PRODUCES	CHECK BY
`NTP × Knowledge`	Hallucinated specifics (citations, quotes, statistics)	Search for the source independently
`Context window × Steerability`	Drift in long conversations (instructions forgotten, tone shifts)	Compare output against original instructions
`Tokenization × NTP`	Confident arithmetic errors	Recalculate independently

Framework: Anthropic 4D AI Fluency Framework, Discernment competency (Dakan, Feller, and Anthropic, 2025).